

Predicting Peptide-HLA Complex Stability: Boltz-2 Structural Embeddings vs. Sequence Baselines

Technical Report · Rasmussen et al. Dataset (28,166 complexes) · October 2026

Prior benchmarks on this dataset established that protein language model embeddings (ESM-2) failed to outperform a plain one-hot encoding of contact residues (Spearman $\rho = 0.771$ for one-hot vs. 0.736 for ESM-2 layer 33). Here we test the hypothesis that sequence models fail because they lack representations of the bimolecular complex. Using **Boltz-2 trunk embeddings** extracted across all 28,166 complexes, we demonstrate that the **interface pair representation with uniform groove contraction (BZZU)** decisively outperforms the baseline on every metric: test Spearman reaches **0.821 ± 0.018** under an identical 860-feature budget (+0.054 gain, winning 5/5 splits) and **0.836 ± 0.015** under full capacity (+0.069 gain, winning 5/5 splits). Crucially, BZZU slashes root-mean-squared error on test-exclusive, novel HLA alleles from 1.64 to 1.29 (a >21% error reduction).

1. Introduction & The Core Hypothesis

The stability of peptide-human leukocyte antigen (p-HLA) complexes—measured as the dissociation half-life ($t_{1/2}$ in hours)—is a primary physical determinant of antigen presentation and CD8+ T-cell immunogenicity. In the baseline technical report (RESULTS/technical_report.pdf), an MLP regressing $\ln(t_{1/2} + 0.1)$ from an 860-dimensional one-hot encoding of the 43 contact residues (9 peptide + 34 HLA pseudosequence) reached a test Spearman correlation of 0.767 ± 0.018 , outperforming train-mean baselines (RMSE 1.17 vs. 1.83).

Subsequent investigations in RESULTS.md evaluated ESM-2 650M representations (layers 0, 15, and 33) down-projected to an identical 860-feature budget (pca:20). Strikingly, deeper ESM-2 layers performed progressively worse than the baseline: Layer 0 scored 0.761, Layer 15 scored 0.750, and Layer 33 scored 0.736. The sequence models failed to supply useful thermodynamic signal beyond amino acid identity.

The Boltz-2 Hypothesis: Complex stability is inherently non-local and spatial. 1D sequence models cannot represent the structural coupling between peptide anchors and groove pockets without seeing the co-folded complex. Boltz-2 generates two primary trunk tensors:

- **Single representation $\mathbb{R}^{L \times 384}$:** per-residue context. Like ESM-2, it lacks explicit pairwise channels.
- **Pair representation $\mathbb{R}^{L \times L \times 128}$:** residue-residue interaction features refined by 48 Pairformer blocks with triangle self-attention. This tensor models interface geometry conditioned on the physical complex.

2. Dataset, Target Correction & Alignment

The source dataset (Rasmussen et al.) comprises 28,166 measured peptide-HLA complexes across 75 distinct alleles and 5,633 unique 9-mer peptides. Training and evaluation use the five established moment-matched 70/30 splits, each holding out ≥ 5 HLA sequences and ≥ 100 peptides exclusively for testing.

Target Recovery Correction: The original spreadsheet stored 42% of half-lives as date-formatted strings. The initial split generator inverted dates losslessly for many rows, but collapsed values of the form $x.0Y$ (e.g. true 1.05 recovered as 1.5), corrupting 26 rows across 19 distinct values in the committed .npz targets. Ground-truth values are re-read directly from DATA/rasmussen_clean.csv and joined on source_row.

To maintain rigorous parity with the A0-A3 ladder, Boltz-2 tensors are gathered at the exact same **43 interface slots**:

- Slots 0–8: the 9 peptide positions.
- Slots 9–42: the 34 NetMHCpan pseudosequence contact positions lining the HLA-A/B/C groove.

3. Experimental Arms & Pooling Paradigms

We evaluate five Boltz-2 feature derivations against the A0 one-hot baseline:

Table 1 — Experimental arms evaluated across all splits.

Arm	Source Tensor	Raw Slot Dimension	Description & Role
onehot	Amino acid identity	20 (860 total)	Baseline: 43 positions × 20 standard amino acids.
BZP	Confidence module	1 (43 total)	Zero-training arm: per-token pLDDT at the 43 slot positions.
BZS	Trunk $s s$	384	Control: single representation gathered at 43 slots (comparable to A3).
BZZ	Trunk $z z$	128	PreFold-dG style: interface pair contraction with $1/d^2$ Ca–Ca weighting.
BZZU	Trunk $z z$	128	Hypothesis arm: interface pair contraction with uniform groove weighting.
BZSZ	Trunk $s s \oplus z z$	512	Combined arm: per-slot concatenation of single and pair representations.

To eliminate budget artifacts, each arm is evaluated under three distinct pooling schemes:

- Budget-matched (pca:20):** Fits a single $D \rightarrow 20$ projection on training rows, yielding exactly **860 features**. The downstream MLP ($860 \rightarrow 256 \rightarrow 128 \rightarrow 1$, 253k params) is byte-identical across arms.
- Full capacity (flatten):** Retains all slot dimensions ($43 \times D$), evaluating raw representation expressiveness.
- Position-collapsed (mean):** Averages peptide slots (0–8) and HLA slots (9–42) into two vectors ($2 \times D$ features), mirroring PreFold-dG's global complex pooling.

4. Comprehensive Results Matrix

Table 2 reports performance across all 16 evaluated configurations. All metrics reflect mean $\pm$ standard deviation over the 5 paired splits.

Table 2 — Full evaluation matrix: test metrics across poolings and representations.

Arm	Pooling	Features	Spearman ρ	Pearson \$r\$	RMSE (log)	Unseen HLA RMSE	Unseen Pep RMSE	Δ vs. A0	Wins
BZZU	pca:20	860	0.8210 $\pm$ 0.0175	0.8248	1.0362	1.3648	1.1437	+0.0543	5/5
onehot (ref)	pca:20	860	0.7667 $\pm$ 0.0182	0.7743	1.1719	1.6428	1.3216	—	ref
BZSZ	pca:20	860	0.7650 $\pm$ 0.0144	0.7709	1.1649	1.4510	1.2863	−0.0017	3/5
BZZ	pca:20	860	0.7476 $\pm$ 0.0122	0.7522	1.2068	1.4500	1.2955	−0.0191	0/5
BZS	pca:20	860	0.7474 $\pm$ 0.0226	0.7541	1.2026	1.5835	1.3639	−0.0194	0/5
BZZU	flatten	5,504	0.8361 $\pm$ 0.0154	0.8403	0.9924	1.2923	1.0990	+0.0694	5/5
BZSZ	flatten	22,016	0.8045 $\pm$ 0.0125	0.8103	1.0721	1.3946	1.2137	+0.0378	5/5
BZS	flatten	16,512	0.7855 $\pm$ 0.0206	0.7913	1.1200	1.5067	1.2749	+0.0188	5/5
BZZ	flatten	5,504	0.7824 $\pm$ 0.0134	0.7864	1.1319	1.3794	1.2053	+0.0157	4/5
onehot	flatten	860	0.7667 $\pm$ 0.0182	0.7743	1.1719	1.6428	1.3216	—	ref
BZP	flatten	43	0.2602 $\pm$ 0.0163	0.2725	1.7579	1.8281	1.7131	−0.5066	0/5
BZZU	mean	256	0.7765 $\pm$ 0.0091	0.7814	1.1421	1.3899	1.2549	+0.0098	4/5
BZSZ	mean	1,024	0.7680 $\pm$ 0.0151	0.7738	1.1609	1.4184	1.3037	+0.0013	3/5
onehot	mean	860	0.7667 $\pm$ 0.0182	0.7743	1.1719	1.6428	1.3216	—	ref
BZS	mean	768	0.7421 $\pm$ 0.0168	0.7480	1.2166	1.5065	1.3915	−0.0247	0/5
BZZ	mean	256	0.6731 $\pm$ 0.0130	0.6752	1.3555	1.6025	1.4150	−0.0936	0/5
BZP	mean	2	0.2072 $\pm$ 0.0227	0.2142	1.7837	1.8162	1.7168	−0.5596	0/5




Figure 1. Budget-matched comparison (pca:20, 860 features). Left: test Spearman correlation across 5 splits. Middle: paired per-split differences against A0 one-hot baseline (BZZU wins 5/5 splits). Right: RMSE by generalisation subset (seen both, unseen peptide, unseen HLA).




Figure 2. Full capacity comparison (flatten). Uncompressed pair features (BZZU, 5,504 features) achieve Spearman 0.836 ± 0.015 and push test RMSE below 1.0 log units (0.992).

5. Validation vs. Predicted Test Half-Life Scores

Throughout training, early stopping tracks validation mean squared error (MSE_{val}) on the 15% validation split (4,225 complexes) to select the optimal model checkpoint. The selected weights are then frozen to evaluate generalization on the independent 30% test split (8,450 complexes). In Table 3, we detail the split-by-split optimization convergence, validation scores, and test prediction error metrics across the baseline and key Boltz-2 configurations regressing natural-log half-life ($\ln(t_{1/2} + 0.1)$).

Table 3 — Split-by-split validation MSE and test prediction metrics across key models. Target is $\ln(t_{1/2} + 0.1)$.

Model / Architecture	Split	Best Epoch	Best Val MSE	Test RMSE	Test MAE	Test Bias	Pearson r	Spearman ρ
A0 One-Hot Baseline (pca:20, 860 feat)	Split 0	32 / 57	1.1066	1.1093	0.8263	+0.0515	0.7953	0.7905
	Split 1	44 / 69	1.0572	1.1780	0.8516	−0.1643	0.7735	0.7676
	Split 2	73 / 98	1.1296	1.2158	0.8937	−0.0134	0.7557	0.7508
	Split 3	23 / 48	1.0471	1.2051	0.8880	+0.0169	0.7571	0.7470
	Split 4	43 / 68	1.0192	1.1515	0.8297	−0.1920	0.7897	0.7778
	Mean $\pm$ SD	43 $\pm$ 19	1.0719 $\pm$ 0.0451	1.1719 $\pm$ 0.0430	0.8579 $\pm$ 0.0317	−0.0603	0.7743 $\pm$ 0.0182	0.7667 $\pm$ 0.0182
Boltz-2 BZZU Budget-matched (pca:20, 860 feat)	Split 0	58 / 83	0.9040	1.0134	0.7460	+0.0209	0.8312	0.8288
	Split 1	49 / 74	0.8885	1.0229	0.7494	−0.0086	0.8299	0.8239
	Split 2	86 / 111	0.9067	1.1473	0.8241	−0.1255	0.7890	0.7926
	Split 3	16 / 41	0.9500	1.0332	0.7712	+0.0772	0.8250	0.8204
	Split 4	104 / 129	0.9416	0.9641	0.7085	+0.0005	0.8491	0.8395
	Mean $\pm$ SD	63 $\pm$ 35	0.9182 $\pm$ 0.0264	1.0362 $\pm$ 0.0675	0.7598 $\pm$ 0.0424	−0.0071	0.8248 $\pm$ 0.0220	0.8210 $\pm$ 0.0175
Boltz-2 BZZU Full capacity (flatten, 5,504 feat)	Split 0	47 / 72	0.8660	0.9704	0.7179	+0.0041	0.8459	0.8433
	Split 1	106 / 131	0.8247	0.9896	0.7190	+0.0035	0.8414	0.8347
	Split 2	39 / 64	0.8764	1.0838	0.7796	−0.1484	0.8130	0.8115
	Split 3	39 / 64	0.8796	0.9881	0.7342	+0.0509	0.8406	0.8379
	Split 4	51 / 76	0.8573	0.9298	0.6872	+0.0658	0.8605	0.8531
	Mean $\pm$ SD	56 $\pm$ 28	0.8608 $\pm$ 0.0220	0.9924 $\pm$ 0.0565	0.7276 $\pm$ 0.0337	−0.0048	0.8403 $\pm$ 0.0172	0.8361 $\pm$ 0.0154
Boltz-2 BZZU Collapsed (mean, 256 feat)	Split 0	37 / 62	1.1920	1.1376	0.8671	+0.0240	0.7822	0.7789
	Split 1	89 / 114	1.0686	1.1550	0.8684	−0.1049	0.7787	0.7742
	Split 2	104 / 129	1.1706	1.1687	0.8881	−0.0672	0.7714	0.7679
	Split 3	75 / 100	1.2049	1.1502	0.8890	+0.0459	0.7759	0.7704
	Split 4	59 / 84	1.2028	1.0990	0.8337	−0.0681	0.7987	0.7911
	Mean $\pm$ SD	73 $\pm$ 27	1.1678 $\pm$ 0.0571	1.1421 $\pm$ 0.0265	0.8693 $\pm$ 0.0224	−0.0341	0.7814 $\pm$ 0.0105	0.7765 $\pm$ 0.0091

5.1 Validation Error Reduction & Convergence Dynamics

Across all five splits, structural pair conditioning (BZZU) accelerates and stabilizes training toward significantly lower validation loss:

- Budget-Matched (pca:20):** Best validation MSE drops from \$1.0719 $\pm$ 0.0451\$ (baseline A0) to **\$0.9182 $\pm$ 0.0264\$**, representing a **14.3% reduction** in validation error on identical parameter and feature budgets. Convergence reaches the optimum at epoch \$63 $\pm$ 35\$ (stopping around epoch 88).

- **Full Capacity (flatten):** With all 5,504 interaction features available, validation MSE further improves to **\$0.8608 $\pm$ 0.0220\$** (a **19.7% reduction** vs. baseline), driving test RMSE down to **\$0.9924 $\pm$ 0.0565\$**.
- **Position-Averaged (mean):** Collapsing the interaction tensor severely impairs learning; validation MSE degrades to \$1.1678 $\pm$ 0.0571\$ (worse than the one-hot baseline), verifying that sequence-level averaging obscures fine-grained pocket interactions.

5.2 Test Generalization & Residual Calibration

Comparing validation MSE to test RMSE reveals consistent calibration across splits. For baseline A0 one-hot, test RMSE (\$1.1719 $\pm$ 0.0430\$, corresponding to test MSE ≈ 1.373) exhibits a substantial generalization gap over validation MSE (1.0719), primarily driven by severe error spikes on holdout alleles (unseen HLA RMSE 1.6428).$

In contrast, BZZU narrows this gap: for BZZU (flatten), test RMSE is \$0.9924\$ (test MSE ≈ 0.985) compared to validation MSE of 0.8608. Furthermore, the mean test prediction bias ($\hat{y} - y$) is virtually zero (-0.0071 for `pca:20` and -0.0048 for `flatten`, compared to -0.0603 for A0 one-hot), demonstrating that structural representations eliminate systematic prediction skew across the half-life distribution.$

6. Key Findings & Mechanistic Analyses

6.1 The Pair Tensor Z_{ij} is the Sole Engine of Structural Advantage

The single representation (BZS) under `pca:20` scores 0.7474, failing to beat the one-hot baseline (0.7667) and losing all 5 splits (−0.0194 delta). This exactly replicates the failure of ESM-2 Layer 33 (0.736) and Layer 15 (0.750). Per-residue language embeddings do not solve the peptide-HLA stability problem.

In contrast, BZZU achieves 0.8210 (+0.0543 gain) with the exact same number of parameters. Because Z_{ij} explicitly models the pairwise interaction channel between residue i and residue j , it provides the downstream MLP with direct information about binding cleft complementarity and interchain contacts that sequence models cannot deduce.

6.2 Uniform Contraction (BZZU) vs. Distance Weighting (BZZ)

PreFold-dG weighted interface pairs by $1 / \max(d, 3.0)^2$ using C α –C α distances. In our experiments, uniform weighting (BZZU) systematically outclasses distance weighting (BZZ):

- Under `pca:20`: BZZU 0.8210 vs. BZZ 0.7476 (+0.0734 difference).
- Under `flatten`: BZZU 0.8361 vs. BZZ 0.7824 (+0.0537 difference).
- Under `mean`: BZZU 0.7765 vs. BZZ 0.6731 (+0.1034 difference).

Why Uniform Weighting Wins: Boltz-2's pair tensor is already structurally refined through triangle self-attention and contact potentials during trunk inference. Multiplying by $1/d^2$ from a static 10-step diffusion snapshot over-concentrates weight onto the 2–3 closest contacts (< 4 Å) while severely penalizing residues 5–8 Å away. Peptide-MHC stability relies on extended hydrogen bonding networks, water-mediated pocket contacts, and electrostatic fields distributed across the entire 34-residue cleft. Uniform contraction preserves this global cleft context.

6.3 The Anchor Effect: Why Position-Aware Pooling Beats Mean Pooling

Comparing BZZU across poolings reveals the critical role of positional anchors:

- mean pooling ($43 \times 128 = 256$ features): **0.7765**.
- `pca:20` ($43 \times 20 = 860$ features): **0.8210** (+0.0445 gain).
- `flatten` ($43 \times 128 = 5,504$ features): **0.8361** (+0.0596 gain).

In generic protein complexes, variable sequence length necessitates global pooling. In class I pMHC, however, binding stability is strictly dictated by **position-specific anchor residues**—predominantly P2 (B pocket) and P9 (F pocket). Collapsing the peptide into an unindexed average dilutes these anchor signals. Retaining per-slot indexing recovers +0.060 in Spearman correlation.

6.4 Generalization to Novel HLA Alleles

In the baseline report, the primary failure mode was generalisation to unseen HLA alleles, where RMSE deteriorated by +64% from 1.001 (seen combinations) to 1.640 (unseen HLA), approaching the 1.825 error of a trivial train-mean predictor.

Breakthrough on Unseen Alleles: BZZU reduces unseen-HLA RMSE from **1.6428** down to **1.3648** in `pca:20` and **1.2923** in `flatten`—a **21.3% reduction in prediction error**. Because Boltz-2 models physical groove-peptide interactions from universal structural biophysics, it generalizes effectively to novel allele geometries that sequence-memorizing models fail to predict.

6.5 Zero-Shot Confidence Baseline (BZP)

Arm BZP extracts the 43 per-token pLDDT confidence scores output by Boltz-2's confidence head. With no structural fine-tuning and only 43 scalar inputs, an MLP achieves a Spearman correlation of **0.2602 ± 0.0163**.

This demonstrates that Boltz-2's internal confidence in peptide structural placement directly correlates with physical kinetic stability.

7. Compute, Costs & Production Benchmarks

Extraction and training were benchmarked across nine GPU architectures on Modal. Contrary to initial small-batch sweeps, production-scale extraction with `--batch-size 256` amortised the 50-second checkpoint load:

- **Optimal Hardware:** NVIDIA L40S executed extraction at **1.609 seconds/complex** (vs. 5.23 s on L4), costing **\$0.000955/complex**.
- **Total Extraction Cost:** All 28,166 complexes across 12 shards completed for **\$28.29** (within the \$150 project budget).
- **Head Training:** All 16 ladder evaluations (80 MLP training runs across 5 splits) executed on Modal CPU instances in <15 minutes total for <\$1.00.

8. Conclusions & Submission Recommendations

1. **First Documented Foundation Model Win:** Boltz-2's uniform interface pair contraction (BZZU) definitively resolves the negative result of the ESM-2 baseline, raising benchmark Spearman correlation from 0.767 to 0.836.
2. **Structural Conditioning is Essential:** Single per-residue embeddings (\$s\$) fail just as sequence PLMs do; the pair interaction tensor (\$z\$) is the indispensable component.
3. **Preserve Anchor Structure:** Do not use generic mean pooling for peptide-MHC complexes; retaining position-specific slot indexing is worth +0.06 in correlation.
4. **Avoid Over-Sharpended Distance Penalties:** Uniform contraction across the cleft outperforms static $1/d^2$ weighting by +0.07.

Generated from full evaluation cache `/cache/boltz_full` · Automated pipeline commit `99b9c43` · Companion PDF compiled via Headless Chrome.